Multiple Linear Regression

Data Exploration

(Multiple) Linear Regression

- In practise there is normally more than one independent variable
- The plane of best fit is given by
-
- $y = a_1 * x_1 + a_2 * x_2 \ldots + a_n * x_n + b$
-
- Again the parameters are obtained by minimising the sum of the squares of the errors.

Multiple Linear regression

$$y = \sum_k a_k x_k + b$$

- y is the dependent variable
- x_k are independent variables
- a_k, b are parameters
- For Example
- StackLoss = a_1*AirFlow + a_2*WaterTemp + a_3*AcidTemp + b

Stackloss Dataset

- https://stat.ethz.ch/R-manual/R-patched/library/datasets/html/stackloss.html
- 21 observations of 4 variables
- Obtained from 21 days of operation of a plant for the oxidation of ammonia (NH3) to nitric acid (HNO3)
- Stackloss (the dependent variable) is 10 times the percentage of the ingoing ammonia to the plant that escapes

Data Exploration

```
import pandas as pd
from pandas.plotting import scatter_matrix
import matplotlib.pyplot as plt


stacklossDF = pd.read_csv("data/stackloss.csv")
# print(stacklossData)

print(stacklossDF.describe())
print(stacklossDF.corr())

scatter_matrix(stacklossDF)
# plt.show()
# plt.savefig('plots/p3stacklossScatter.png')
```

Data Exploration

- pd.DataFrame.describe()
  - Summary of numeric features
  - Mean, Min, Max, std etc.
- pd.DataFrame.corr()
  - Matrix of correlation coefficients between the variables
- pd.scatter_matrix()
  - A matrix of scatterplots

DataFrame.describe()

	AirFlow	WaterTemp	AcidConc	StackLoss
count	21.000000	21.000000	21.000000	21.000000
mean	60.428571	21.095238	86.285714	17.523810
std	9.168268	3.160771	5.358571	10.171623
min	50.000000	17.000000	72.000000	7.000000
25%	56.000000	18.000000	82.000000	11.000000
50%	58.000000	20.000000	87.000000	15.000000
75%	62.000000	24.000000	89.000000	19.000000
max	80.000000	27.000000	93.000000	42.000000

DataFrtame.corr()

	AirFlow	WaterTemp	AcidConc	StackLoss
AirFlow	1.000000	0.781852	0.500143	0.919663
WaterTemp	0.781852	1.000000	0.390940	0.875504
AcidConc	0.500143	0.390940	1.000000	0.399830
StackLoss	0.919663	0.875504	0.399830	1.000000

pandas.plotting.scatter_matrix()

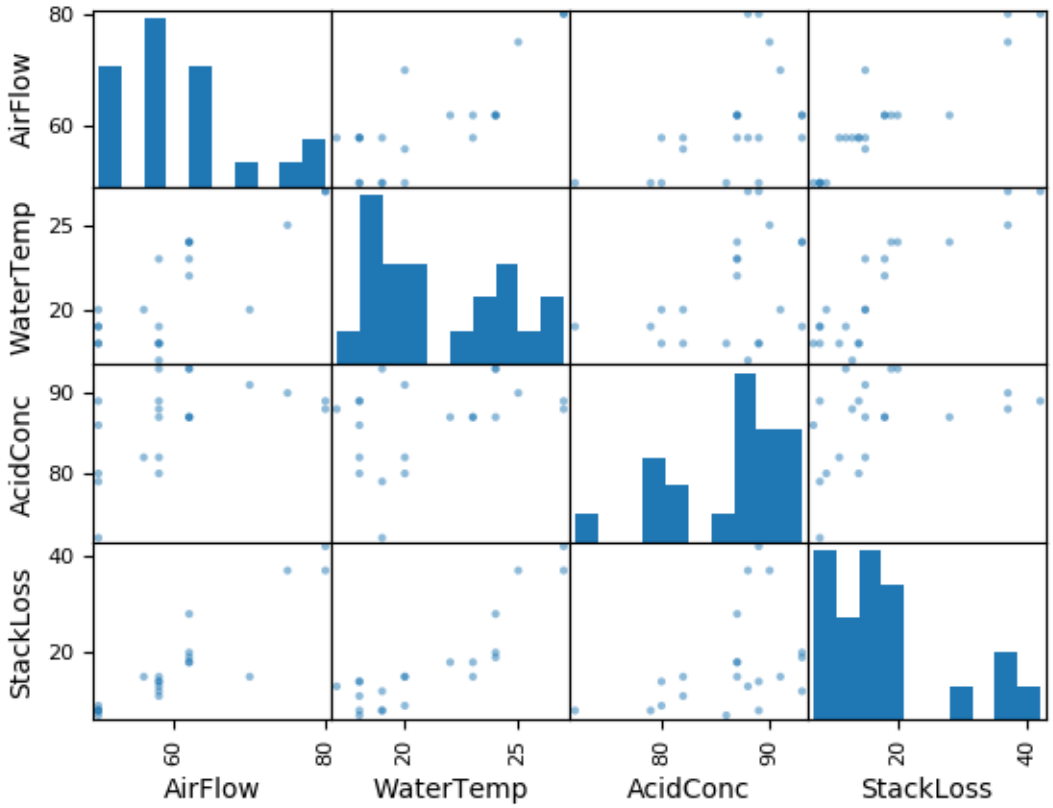

Scatter Plot Matrix - Analysis

- There seems to be a fairly strong linear correlation between StackLoss and AirFlow
- Also between StackLoss and WaterTemp
- Not so strong between StackLoss and AcidConc
- Maybe even a non-linear relationship between StackLoss and AcidConc.
- Even so, the normal thing to do is proceed using all three variables.
- We would not be surprised if AcidConc is not as strong a predictor as the other variables.

Marginal Relationships

- The scatter plots show the marginal relationships between variables without regard to other variables.
- Note that the absence of a correlation between an independent variable (predictor) and the dependent variable (output) does not mean that the dependent variable is not useful as a predictor.
- (For example, points that should be near a line of best fit could be being moved away from it by the values of other variables.)

Marginal Relationships

- With multiple linear regression we are primarily concerned with how the dependent or output variable relates to the independent (predictor) variable simultaneously.

Feature Reduction

- It is nearly always safest to use all the features unless there is a good reason not to.
- If a predictor variable has close to zero variance (it does not vary much) then it can have very little predictor value.
- If two predictor variables are very strongly correlated, then one is redundant and possibly can be omitted.
- For performance reasons a model with a reduced set of features with similar accuracy to the full model can be used.